

From Edge AI to Governed Autonomous Edge Intelligence




Aridio Silva · 11 min read · 15 hours ago



Why autonomous AI agents require security, governance, and bounded authority by design

Aridio Silva

Independent Researcher | Creator of SGAEIA



From Edge AI to Governed Autonomous Edge Intelligence. SGAEIA explores how distributed autonomous AI can combine intelligence and operational authority with Zero Trust, continuous governance, security enforcement, and auditable evidence.

Artificial Intelligence is undergoing a fundamental architectural transition.

For years, AI systems were primarily designed to analyze data, generate predictions, classify information, or produce content. Even with the emergence of generative AI and Large Language Models, most systems

initially remained essentially advisory: they generated information that humans could evaluate before taking action.

Agentic AI changes this relationship.

Autonomous AI agents can increasingly perceive their environments, reason about goals, communicate with other agents, invoke tools, access external services, delegate tasks, and perform actions across digital and physical infrastructures. [3], [5]

At the same time, AI computation is moving beyond centralized cloud environments toward increasingly distributed architectures involving edge devices, gateways, local computing nodes, private infrastructure, and cloud services.

These two transformations — **autonomous agency and distributed intelligence** — create an architectural problem that cannot be addressed simply by adding conventional security controls around an AI model.

The central question becomes:

How can an autonomous AI system be granted enough authority to perform useful actions while ensuring that such authority remains bounded, verifiable, auditable, and revocable?

This question is at the center of my current research and motivated the development of **SGAEIA — Secure Governed Autonomous Edge Intelligence Architecture**. [15]

. . .

From Cloud AI to Edge AI

Cloud computing enabled organizations to centralize massive computational resources and datasets, making the development and deployment of large-scale AI systems economically viable.

However, not every intelligent decision can depend entirely on distant cloud infrastructure.

Latency, connectivity, privacy, resilience, bandwidth, regulatory requirements, operational autonomy, and real-time decision-making increasingly require computation to occur closer to where data is generated.

This has contributed to the development of **Edge AI**.

In a conventional Edge AI architecture, models may execute on devices, gateways, industrial systems, vehicles, local servers, or regional infrastructure. Some processing occurs locally, while more computationally intensive tasks may still be delegated to cloud services. [1], [2]

Edge AI therefore distributes **computation**.

Agentic AI introduces something more consequential.

It also distributes **authority**.

Edge AI distributes computation. Agentic Edge AI distributes computation and authority.




Figure 1 — From Computation to Authority. Traditional Edge AI primarily distributes computation; Agentic Edge AI additionally distributes operational authority, introducing new requirements for security, governance, delegation, and accountability.

This distinction changes the security model.

Protecting models, data, networks, and devices remains necessary, but it is no longer sufficient.

The architecture must also control **what autonomous agents are authorized to do**.

. . .

When Intelligence Becomes Authority

Consider the difference between two AI systems.

The first analyzes telemetry from an industrial environment and recommends that a machine should be stopped.

The second analyzes the same telemetry, determines that the machine represents an operational risk, authenticates itself to an industrial control system, invokes an authorized tool, stops the machine, creates an incident record, and delegates diagnostic analysis to another autonomous agent.

Both systems use AI.

Architecturally, however, they belong to very different risk categories.

The first system produces **information**.

The second exercises **authority**.

Once an AI agent can invoke tools, modify systems, communicate with other agents, access resources, delegate tasks, or trigger real-world actions, security can no longer focus exclusively on protecting models and data. [3], [5]

We must also govern the **authority exercised by the agent**.

. . .

The Problem of Autonomous Delegation

Multi-agent systems make this challenge substantially more complex.

An agent may receive authority from a human, application, organization, or another agent. It may then delegate part of a task to another agent, which may invoke another service or agent.

A relatively simple workflow can therefore evolve into a chain such as:

Human → Agent A → Agent B → Tool / Service → External System

Authentication alone cannot govern this chain.

The architecture also needs to determine what authority was originally granted, which portion may be delegated, to whom, for what purpose, for how long, under which contextual conditions, and whether that authority can subsequently be reduced or revoked. [4], [8]

Most importantly, the complete chain must remain reconstructable.




Figure 2 — Chain of Delegated Authority in a Multi-Agent System. Authority originates from an explicit grant and may propagate through agents, tools, and external systems only within enforceable constraints, while remaining observable, auditable, and revocable.

This means identity alone is insufficient.

Authentication answers:

Who are you?

Governed autonomous systems must additionally answer:

What are you allowed to do, why are you allowed to do it, under which conditions, and who granted that authority?

Delegation must therefore never imply uncontrolled expansion of authority.

An agent receiving a subset of authority should not be able to silently create a larger authority for another agent.

This suggests a fundamental architectural invariant:

Delegated authority must remain equal to or narrower than the authority from which it originated. [15]

. . .

Bounded and Revocable Authority

One of the central principles explored in SGAEIA is that autonomous agents should not receive unrestricted or implicitly permanent authority.

Agent authority should be **bounded**.

A useful conceptual representation is:

Authority = Identity + Scope + Resource + Action + Context + Time + Delegation Constraints [15]

Each component answers a different security question.

Identity establishes who is requesting authority.

Scope determines the domain within which that authority is valid.

Resource identifies what may be accessed.

Action specifies which operations may be performed.

Context constrains authority according to environmental, operational, or risk conditions.

Time establishes when and for how long authority remains valid.

Delegation Constraints determine whether authority can be transferred and how far that delegation may propagate.




Figure 3 — Bounded and Revocable Authority. Agent authority is modeled as the combination of identity, scope, resource, action, context, time, and delegation constraints, continuously subject to policy enforcement, monitoring, audit, and revocation.

Authority must also be **revocable**.

If an agent becomes compromised, behaves unexpectedly, violates policy, or operates under conditions that are no longer valid, the system must be capable of reducing or terminating its authority.

Revocation should not depend solely on stopping the AI model.

The enforcement infrastructure itself must be capable of rejecting future actions. [9]

This distinction is fundamental.

A trustworthy architecture cannot rely on an autonomous agent voluntarily respecting the limits imposed upon it. Those limits must be externally enforceable.

. . .

Zero Trust for Agentic AI

This requirement leads naturally to **Zero Trust**.

Traditional security architectures often assume that entities operating within trusted network boundaries deserve some degree of implicit trust. [9]

Distributed autonomous systems make this assumption increasingly dangerous.

An AI agent may execute on an edge device, migrate between environments, invoke cloud services, collaborate with remote agents, access external tools, or operate across administrative domains.

Physical or network location therefore provides insufficient evidence of trustworthiness. [9], [10]

A Zero-Trust approach to Agentic AI implies that actions should be continuously evaluated according to identity, authorization, policy, context, resource sensitivity, delegation state, and other relevant evidence. [9]

The principle becomes:

**Never trust an agent merely because of where it is executing.
Verify whether the requested action is authorized at the moment
it is requested.**

Security consequently moves from a static perimeter toward a **continuous runtime decision process**.

. . .

Governance Must Become Executable

Governance is frequently expressed through policies, standards, procedures, spreadsheets, documents, and compliance reports.

These mechanisms remain important.

But autonomous systems operate at machine speed.

A policy that exists only in a document cannot directly prevent an unauthorized autonomous action.

For increasingly autonomous infrastructures, governance must therefore become more **machine-readable, enforceable, observable, and testable**. [8], [11]

This leads to another principle explored in SGAEIA:


Governance should progressively become an executable architectural property. [15]

Policies define constraints.

Runtime components enforce them.

Telemetry records decisions.

Evidence demonstrates whether controls were applied.



Tests verify whether security invariants remain valid.

Governance, security, observability, and assurance can therefore become components of the same engineering lifecycle.

. . .

Evidence-as-Code and Continuous Assurance

Traditional compliance frequently evaluates systems periodically.

Autonomous AI systems can change their behavior and operational context much faster than traditional audit cycles. [11]

This creates the need for **continuous assurance**.

SGAEIA explores the concept of **Evidence-as-Code**, in which evidence concerning architecture, policy enforcement, tests, security controls, and system behavior can be generated and evaluated through automated processes. [15]

The objective is not to eliminate human governance.

It is to provide humans with **stronger, more timely, and verifiable evidence**.

Instead of asking only:

Was this system compliant during the last audit?

we should increasingly be able to ask:

What evidence demonstrates that this autonomous system is operating within its authorized boundaries now?

. . .

Security-by-Design for Autonomous Systems

Many security problems become extremely difficult to correct after an architecture has already been deployed.

This is particularly true for distributed multi-agent systems.

If identity, authority, delegation, revocation, auditability, trust boundaries, and policy enforcement are not architectural concerns from the beginning, adding them later may require fundamental redesign.

For this reason, SGAEIA combines principles associated with **Security-by-Design, Security-First, and Shift-Left engineering**. [12], [14], [15]

Security-by-Design treats security properties as architectural requirements.

Security-First ensures that security assumptions influence engineering decisions from the beginning.

Shift-Left introduces threats, security requirements, invariants, and tests early in the development lifecycle. [12], [13], [14]

Threat modeling approaches such as **STRIDE**, together with applicable **OWASP** guidance, can then be integrated with specifications, tests, runtime enforcement, and evidence generation. [6], [7], [13]

. . .

Toward Governed Autonomous Edge Intelligence

Edge AI traditionally focuses on moving intelligence closer to data sources and operational environments. [1], [2]

Agentic AI introduces autonomous decision-making and action. [3], [5]

When these two paradigms converge, a new architectural category begins to emerge.

I refer to this direction as:

Governed Autonomous Edge Intelligence

The distinction is important.

Distributed intelligence alone is not sufficient.

Autonomy alone is not sufficient.

A sustainable architecture for distributed autonomous AI must combine:

Intelligence + Authority + Security + Governance + Evidence

The goal is not to prevent autonomous systems from exercising meaningful authority.

The goal is to make that authority **explicit, bounded, traceable, enforceable, auditable, and revocable.**

. . .

SGAEIA

SGAEIA — Secure Governed Autonomous Edge Intelligence Architecture is an open and technology-neutral research initiative exploring these architectural requirements. [15]

Rather than defining a single technology stack, SGAEIA investigates architectural contracts, security invariants, governance mechanisms, and assurance properties that should remain valid even when underlying technologies change.




Figure 4 — Governed Autonomous Edge Intelligence (SGAEIA). Intelligence, bounded authority, security, governance, and evidence converge into an architectural model for trustworthy autonomous multi-agent operations across heterogeneous edge-cloud environments.

The architecture brings together five fundamental dimensions:

Intelligence — perception, reasoning, planning, decision-making, and collaborative intelligence.

Authority — explicit, bounded, delegated, enforceable, and revocable permission to act.

Security — identity, authentication, Zero Trust, policy enforcement, threat detection, and runtime protection.

Governance — machine-readable policies, risk management, continuous GRC, accountability, compliance, and traceability.

Evidence — observability, auditability, Evidence-as-Code, continuous assurance, and artifacts capable of supporting verification.

These dimensions operate across distributed multi-agent systems spanning edge devices, industrial systems, vehicles, IoT sensors, gateways, robotics, private infrastructure, regional nodes, and cloud services.

The objective is not merely autonomous operation.

It is **governed autonomous operation**.

. . .

Security-Preserving Technology Substitution

The Agentic AI ecosystem is evolving extremely rapidly.

Models, frameworks, agent runtimes, communication protocols, vector databases, identity systems, policy engines, orchestration technologies, and infrastructure components will inevitably change.

An architecture that depends on a specific implementation risks becoming obsolete as those technologies evolve.

SGAEIA therefore explores another important principle: **security-preserving technology substitution**. [15]

A component should be replaceable without silently weakening the architectural security properties associated with it.

In other words:

Technologies may change. Security invariants should remain enforceable.

This is one reason SGAEIA is conceived as a **technology-neutral reference architecture**, rather than as a prescription for one particular software stack.

. . .

A Research Question for the Agentic Era

The evolution from generative AI toward autonomous agents fundamentally changes the security problem. [3], [5]

The question is no longer only:

Can the AI generate the correct answer?

Increasingly, we must also ask:

Should this agent be allowed to perform this action?

Who granted that authority?

Can that authority be delegated?

Under what conditions?

Can it be revoked?

Can we prove what happened afterward?

These questions sit at the intersection of Artificial Intelligence, cybersecurity, distributed systems, software architecture, governance, and software engineering.

As AI agents become more autonomous and increasingly distributed across edge-cloud environments, answering these questions will become essential to building trustworthy AI systems.

The transition can therefore be summarized in a simple architectural proposition:

Edge AI distributes computation. Agentic Edge AI distributes computation and authority. Governed Autonomous Edge Intelligence makes that authority explicit, bounded, enforceable, auditable, and revocable.

And this leads to the principle that summarizes the direction of SGAEIA:

Autonomous AI. Governed by Design. Trusted by Evidence.

. . .

References

[1] Zhou, Z., Chen, X., Li, E., Zeng, L., Luo, K., & Zhang, J. (2019). *Edge Intelligence: Paving the Last Mile of Artificial Intelligence With Edge Computing*. Proceedings of the IEEE, 107(8), 1738–1762.

DOI: **10.1109/JPROC.2019.2918951**.

[2] Singh, R., & Gill, S. S. (2023). *Edge AI: A Survey*. Internet of Things and Cyber-Physical Systems, 3, 71–92. DOI: **10.1016/j.iotcps.2023.02.004**.

[3] Wang, L., Ma, C., Feng, X., et al. (2024). *A Survey on Large Language Model Based Autonomous Agents*. Frontiers of Computer Science, 18, 186345. DOI: 10.1007/s11704-024-40231-1.

[4] South, T., Marro, S., Hardjono, T., Mahari, R., Whitney, C. D., Chan, A., & Pentland, A. (2025). *Position: AI Agents Need Authenticated Delegation*. Proceedings of the 42nd International Conference on Machine Learning, Proceedings of Machine Learning Research, 267, 82211–82231.

[5] Chhabra, A., Datta, S., Nahin, S. K., & Mohapatra, P. (2025). *Agentic AI Security: Threats, Defenses, Evaluation, and Open Challenges*.

arXiv:2510.23883. DOI: 10.1007/s11704-024-40231-1.

[6] OWASP GenAI Security Project. (2025). *Agentic AI — Threats and Mitigations*. OWASP Agentic Security Initiative.
OWASP — Agentic AI Threats and Mitigations

[7] OWASP GenAI Security Project. (2026). *OWASP Top 10 for Agentic Applications 2026*.
OWASP Top 10 for Agentic Applications 2026

[8] OWASP GenAI Security Project. (2026). *Agent Control Standard (ACS)*.
OWASP Agent Control Standard

[9] Rose, S., Borchert, O., Mitchell, S., & Connelly, S. (2020). *Zero Trust Architecture*. NIST Special Publication 800–207.
DOI: **10.6028/NIST.SP.800–207**.

[10] Chandramouli, R., & Butcher, Z. (2023). *A Zero Trust Architecture Model for Access Control in Cloud-Native Applications in Multi-Cloud Environments*. NIST Special Publication 800–207A.
DOI: **10.6028/NIST.SP.800–207A**.

[11] Tabassi, E. (2023). *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. NIST AI 100–1. DOI: **10.6028/NIST.AI.100–1**.

[12] Scarfone, K., Souppaya, M., & Dodson, D. (2022). *Secure Software Development Framework (SSDF) Version 1.1: Recommendations for Mitigating the Risk of Software Vulnerabilities*. NIST Special Publication 800–218. DOI: **10.6028/NIST.SP.800–218**.

[13] OWASP Foundation. *Threat Modeling Cheat Sheet*. OWASP Cheat Sheet Series. OWASP Threat Modeling Cheat Sheet

[14] Cybersecurity and Infrastructure Security Agency (CISA) et al. (2023). *Shifting the Balance of Cybersecurity Risk: Principles and Approaches for Secure by Design Software*.

[15] Silva, Aridio. (2026). *SGAEIA — Secure Governed Autonomous Edge Intelligence Architecture*. Research Artifact. Zenodo.
DOI: **10.5281/zenodo.22557796**.

. . .

About the Author

Aridio Silva is an independent researcher based in Brazil working on the architecture, security, governance, and trustworthiness of autonomous and distributed Artificial Intelligence systems.

His research focuses on **Agentic AI, Multi-Agent Systems, Edge AI, AI Security, Zero Trust, Security-by-Design, AI Governance, Spec-Driven Development, and continuous security assurance.**

He is the creator and lead researcher of **SGAEIA — Secure Governed Autonomous Edge Intelligence Architecture**, an open research initiative investigating architectural foundations for secure, governed, auditable, and trustworthy autonomous AI systems operating across distributed edge-cloud environments.

Research & Project Resources

[SGAEIA — GitHub Repository](#)

[SGAEIA — Zenodo Research Artifact / DOI](#)

[ORCID — Aridio Silva](#)

[Google Scholar — Aridio Silva](#)

[GitHub — Aridio Silva](#)

[LinkedIn — Aridio Silva](#)

Figures

Conceptual diagrams developed for **SGAEIA — Secure Governed Autonomous Edge Intelligence Architecture**. © 2026 Aridio Silva. See the SGAEIA repository for applicable licensing and reuse terms.

Artificial Intelligence



Cybersecurity



Edge Computing



Agentic Ai





Written by Aridio Silva

Follow

0 followers · 1 following

Independent Researcher in Agentic AI, Edge AI, Multi-Agent Systems, AI Security & Governance | Creator of SGAEIA